What v0 shows and what it does not

Operate the pipeline without human judgment (tickets)	**shown**	A: 44 tickets, 5 models; E: 6 tickets, 0 owner decisions
Never waste compute on a bad resubmission	**shown**	A: 0 burning answers in the harness arm, 4–7 whiteboard
Lift the cheapest model to frontier level on operations	**shown**	A: Haiku 36.25 > Fable whiteboard 28.5
Choose Wannier projections (the expert step)	**shown**	B: code 13/16 round 1; frontier models 9–12/16
Reproduce a sealed answer end to end, unattended	**shown**	E: GaP, NbFeSb identical; NbFeSb = production flight
Do it cheaper than a frontier model working alone	**shown**	E: 3.6× less compute, 30× less model spend
Pick the next batch of materials better than chance	**partly shown**	C: hits 4–5 vs random 3.1; pool is prior-biased, n = 23
Be faster in wall-clock	**not shown / gap**	E: 9–11 h vs 4–10 h; the queue policy cost 7.5 h
Predict the kernel peak from cheap data	**weak**	D: +1–3 materials over the model alone; locator weak
Solve the hardest Wannier cases with a model (NbFeSb)	**weak**	B: no model passed NbFeSb from the card; harness used the casebook
Guard against silent input differences (cell)	**not shown / gap**	E: SnS2 fell back to a PBEsol cell with one log line
Generalise beyond this pipeline / these 16–44 items	**not shown / gap**	one harness, one lab, n = 3 repeats; no external re-run yet